

Call for Expressions of Interest

<i>Date of Issue:</i>	September 2026
<i>Expected Completion Date:</i>	May 2027
<i>Project Title:</i>	Legal Pythia: SaaS Compliance Intelligence Platform for Fintech and Journalism
<i>Company:</i>	Legal-Pythia LLP
Contact Details	
<i>Lead Contact:</i>	Jeremy Bormann
<i>Lead E-mail:</i>	jeremy@legalpythia.co.uk
<i>Additional Contact:</i>	N/A
<i>E-mail:</i>	N/A
<i>Additional Contact:</i>	N/A
<i>E-mail:</i>	N/A
<i>Additional Contact:</i>	N/A
<i>E-mail:</i>	N/A
Additional Info	
<i>Students must sign NDA:</i>	X
<i>Project uses personal data:</i>	No
<i>Custom hardware required:</i>	No, just laptop or PC If Applicable: Please provide details
<i>Relevant research areas (if applicable)¹:</i>	N/A
Project Background Information	
Legal Pythia LLP is a legal technology company based in Germany, developing explainable AI (XAI) solutions for compliance, fraud detection, and regulatory risk management across European markets. The company's core mission is to make complex legal and regulatory reasoning transparent, auditable, and accessible to non-technical professionals through intelligent software tools. This project proposes the development of a SaaS platform — the Legal Pythia Compliance Intelligence Platform — that serves two high-demand verticals: fintech and journalism. Both sectors face rapidly evolving regulatory landscapes and significant liability exposure, yet lack affordable, explainable AI tooling tailored to their specific workflows. In the fintech configuration, the platform ingests transaction records, customer profile data, and relevant regulatory texts (e.g. AML/CTF directives, KYC requirements, MiCA, and EU consumer protection law) to surface explainable risk indicators, flag anomalies, and generate human-readable compliance rationales. These outputs are designed to support compliance officers in	

¹ An overview of different research areas in the school of computer science can be found [here](#).

making auditable, defensible decisions without requiring deep technical expertise.

In the journalism configuration, the platform assists editors, reporters, and legal teams by analysing draft articles, source documents, and applicable press codes or platform policies. It highlights potential defamation risks, conflicts with journalistic ethics standards, and misinformation indicators — again with clear, rule-grounded explanations that support editorial decision-making.

Across both verticals, the system is designed with explainability and trust at its core. Every output includes a rationale traceable to specific data inputs and regulatory rules, enabling users to understand, challenge, and act on AI-generated recommendations. The platform is envisioned as a multi-tenant SaaS product with role-based access, API integrations, and a modular architecture that allows deployment across different regulatory jurisdictions.

The goal of this student project is to design and build a functional prototype of this platform, validating core technical components (data ingestion, XAI model integration, explainability outputs, and front-end UX) against realistic use case scenarios in at least one of the two target verticals.

Requirements

The group must provide:

A deployable web application prototype with full source code repository (GitHub or equivalent).

Technical documentation covering system architecture, data model, XAI model selection, and data protection considerations.

A data ingestion pipeline handling both structured data (e.g. transactions, article metadata) and unstructured text (regulations, policies, articles).

At least one integrated explainable AI model for risk scoring, anomaly detection, or rule-based compliance checking, with human-readable output rationales.

A usable front-end interface suitable for non-technical professionals (compliance officers, editors, or journalists), with role-based access (e.g. analyst vs reviewer).

Appropriate handling of personal and sensitive data using anonymisation or synthetic datasets throughout development and testing.

A short evaluation report or presentation demonstrating the prototype on at least two realistic scenarios in the fintech or journalism vertical, including sample data and user feedback or expert review.

Skills	
<i>Highly desirable:</i> Machine learning or data science, particularly explainable AI (XAI) or model interpretability Python and common ML/NLP libraries (e.g. scikit-learn, PyTorch, spaCy, Hugging Face) Full-stack or front-end web development (e.g. React, Vue, FastAPI, or Django) Familiarity with EU regulatory frameworks relevant to fintech (AML, KYC, MiCA) or media (press codes, defamation law, GDPR) Database design and querying (SQL and/or NoSQL)	<i>Desirable:</i> Experience with REST APIs and microservice or event-driven architectures UX/UI design skills, especially for decision-support or dashboard tools Familiarity with legal tech, fintech compliance, or digital journalism workflows Experience with cloud deployment (e.g. AWS, GCP, Azure, or Docker/Kubernetes) Knowledge of GDPR and data anonymisation techniques Interest in explainable AI ethics, algorithmic accountability, or regulatory technology (RegTech)
Submission / Award information	
<i>Date of Submission:</i>	September 2026
<i>Date of Pitch:</i>	October 2026
<i>Notification of award:</i>	October 2026

For further information on this proposal, the team representative should in the first instance e-mail the lead contact stated above at the e-mail address provided.